

APOORVA ANURAG

apoorvagk27@gmail.com • 8299133680 • Bengaluru, India • [linkedin.com/in/apoorvaanurag](https://www.linkedin.com/in/apoorvaanurag)

PROFESSIONAL SUMMARY

Data Engineer II with 3 years building production-grade distributed data pipelines in a fintech NBFC — spanning ETL/ELT design, datalake architecture, entity resolution at 13M+ records, and LLM-powered ingestion workflows. Expert in Python, PySpark, SQL, AWS, Databricks, and Apache Airflow, with a strong foundation in data modeling, schema design, performance tuning, and data quality across lending and compliance domains.

TECHNICAL SKILLS

Languages & Processing: Python (OOP), PySpark, SQL, Shell Scripting

Big Data & Distributed Computing: Apache Spark, Databricks, EMR, Dataproc, Distributed Data Processing

Cloud Platforms: AWS (S3, EC2, Redshift, EMR), GCP-equivalent (BigQuery/Redshift, Cloud Storage/S3, Dataproc/Databricks)

Data Warehousing: Redshift, Snowflake-compatible schema design, Data Modeling, SCD Type 2, Star/Snowflake Schema

Workflow Orchestration: Apache Airflow, ETL/ELT Pipeline Design, Workflow Scheduling

DevOps & Quality: Git, Jenkins, Docker, CI/CD Pipelines, Data Quality Frameworks, Unit Testing, Audit Logging

AI / LLM: GPT-4, OpenAI API, LangChain, ChromaDB, RAG Pipeline Design

WORK EXPERIENCE

Senior Associate – Data Engineering Jul 2023 – May 2026 Vivriti Capital Limited • Bengaluru, India

- **Colending Data Platform:** Architected end-to-end PySpark pipelines on Databricks processing 30,000+ borrower records daily across a Rs.10,000Cr AUM portfolio. Replaced manual, vendor-dependent Excel workflows with a unified SQL-queryable analytics layer for Finance and Risk teams, unlocking Rs.80L+ in monthly selldown value across 2 lending pools.
- **Data Modeling & Schema Design:** Designed AWS datalake with SCD Type 2 historization for full point-in-time reconstructability. Implemented star schema data modeling and dimensional schema design to support reliable lending and compliance reporting.
- **Data Validation Framework:** Built a PySpark validation framework with 10+ checks for the weekly loanbook pipeline — enforcing loan ID uniqueness across monthly cycles, schema integrity, selldown date correctness, and partner-level disbursement reconciliation against staging tables. Failures triggered automated Microsoft Teams alerts; framework caught a critical schema mismatch in selldown files that had dual-tagged loans across two pools, preventing incorrect loanbook publication.
- **UCIC Borrower Identification System:** Built an RBI-compliant entity resolution engine in PySpark deduplicating 13M+ borrower records via deterministic and fuzzy matching. Embedded automated discrepancy flagging for upstream data quality assurance and audit readiness. Tuned Spark execution plans and partition strategies to handle record-scale distributed processing efficiently.

- **OnePOS Reconciliation Engine:** Automated reconciliation of core banking ledger against bank statements using optimized SQL queries with performance tuning for large dataset joins and aggregations. Eliminated manual book and account closure workflows, saving 72+ man-hours per cycle.
- **LLM-Powered Ingestion Pipeline:** Designed and deployed a GPT-powered schema standardization pipeline for 30+ monthly collection reports with inconsistent structures. Fed only top-10 rows per file to minimize API latency and cost; LLM returned strict JSON column mappings with confidence scores. Implemented auto-flagging for low-confidence outputs, full decision audit logging, and downstream datalake ingestion. Sole engineer end-to-end.
- **Workflow Orchestration & DevOps:** Orchestrated all batch workflows using Apache Airflow DAGs with dependency management, retry logic, and SLA monitoring. Applied CI/CD pipelines and Git-based version control with code review and testing practices across all production pipelines.
- **Cross-Functional Collaboration:** Partnered with Finance, Risk, and Operations stakeholders to translate business data requirements into scalable pipeline designs, reducing reporting turnaround and enabling faster cross-functional decision-making.

Quantitative Analyst Intern *Jun 2022 – Aug 2022* Investmint • Bengaluru, India

- Built a Python backtesting framework for equities, fixed income, and commodities, executing 1,000+ backtests across varied parameters. Applied walk-forward analysis to optimize trading models, yielding a 12% improvement in Sharpe ratio.

PROJECTS

Credit Document Intelligence Pipeline | Python, OpenAI, LangChain, ChromaDB, RAG

- Built a production-grade RAG pipeline ingesting RBI circulars and credit policy PDFs. Chunks and embeds documents via OpenAI (text-embedding-3-small), stores vectors in ChromaDB, and answers natural language queries with source citations. Designed for scalable document intelligence in lending compliance use cases.

EDUCATION

B.Tech – Computers and Communication Engineering *Jul 2019 – Jun 2023* Manipal Institute of Technology • Karnataka
CGPA: 7.55

CERTIFICATIONS

- AWS Certified Solutions Architect – Associate